


**When we get sick,
we go to a doctor. 🩺**

**But what happens when
a business gets sick?**

**Who diagnoses what's wrong?
Who figures out what needs to be done?
And who actually takes action?**



meet

OpsDoctor

Your AI doctor for business operations.



Diagnose. Decide. Act.

Turning scattered operational symptoms into actionable outcomes.



Business incidents don't live in one system.


They show up as **scattered symptoms across the business.**



Humans spend **valuable time** stitching these signals together before they can even start making a decision.

Meet OpsDoctor

An AI agent that investigates the problem, not just the question.



Reason → Act → Observe → Reason again

The agent dynamically decides what to investigate, which tools to use, and what action to take.

This isn't an API wrapper. It's an agent.

TRADITIONAL AUTOMATION

- Fixed workflow
- Predefined API sequence
- One system at a time
- Limited context
- No dynamic investigation

OPSDOCTOR

- ✓ Understands natural language
- ✓ Chooses tools dynamically
- ✓ Investigates across systems
- ✓ Correlates evidence
- ✓ Evaluates policies
- ✓ Stages actions for approval

The agent determines what needs to happen next.

One agent. Six operational systems. One execution layer.

OpsDoctor connects to your tools, understands the context, and takes action — all through one agent.



OpsDoctor can answer operational questions instantly.

“Which payment gateway is experiencing more failures?”

PAYPAL

8 transactions
4 failed
50% failure rate
\$11,780 volume

STRIPE

5 transactions
1 failed
20% failure rate
\$4,100 volume

PayPal shows the higher observed failure rate: 50% vs 20%.

Connected demo dataset — not a real-world performance claim.

Autonomy with a safety boundary.

AI AGENT

Investigates
Correlates
Evaluates
Recommends
Prepares action

HUMAN


Reviews
Approves or rejects
Controls execution

SWYTCHCODE

Receives approval
Executes operation
Returns result

Reason autonomously — never act consequentially in silence.

An operations command center, not a chatbot.



Turn operational investigation into an agentic workflow.

BEFORE / MANUAL

Alert

Open systems

Search tickets

Find policy

Correlate manually

Write update

Escalate

AFTER / AGENTIC

Ask OpsDoctor

Agent investigates

Evidence correlated

Policy evaluated

Action prepared

Human approves

Execute

Less context switching. Faster investigation. Traceable decisions.

Built to demonstrate real agentic execution.

APPLICATION — OpsDoctor web interface

AGENT LAYER — Intent • Dynamic tools • Multi-step investigation • Evidence synthesis

CORE — AI reasoning • Autonomous ReAct-style planning


EXECUTION — Swytchcode

INTEGRATIONS — PayPal • Stripe • Jira • Slack • Gmail • Notion

The agent is the orchestrator. Swytchcode is the execution layer.

Maru Before:

The business is sick, but diagnosis is manual.



Maru Before:

After: The business gets a diagnosis and a treatment path.



Signals scattered across 6 systems
Humans hunt for relevant evidence
Context lives in different tools
Teams manually correlate incidents
Decisions happen after investigation
Actions require separate handoffs

One natural-language request
Agent selects relevant tools
Evidence correlated across systems
Policies evaluated automatically
Action proposed with context
Human approves → Swytchcode executes

NOW

Live Demo

01 Ask which gateway has more failures

02 Agent investigates payment systems

03 Investigate the PayPal incident

04 Correlate six connected systems

05 Evaluate policy threshold

06 Stage the action

07 Human approves

08 Swytchcode executes

OpsDoctor

From operational noise to actionable decisions.

UNDERSTAND
Natural-language
business problems

INVESTIGATE
Cross-system
autonomous reasoning

ACT
Controlled execution
through Swytchcode

AI that doesn't just answer. It investigates, decides what to do next, and acts with control.